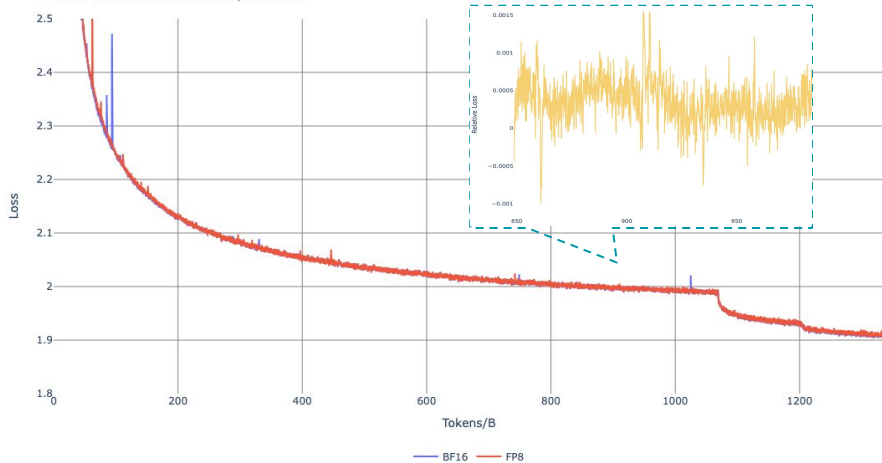


BF16 v.s. FP8 on 16B DeepSeek-V2



BF16 v.s. FP8 on 230B DeepSeek-V2

